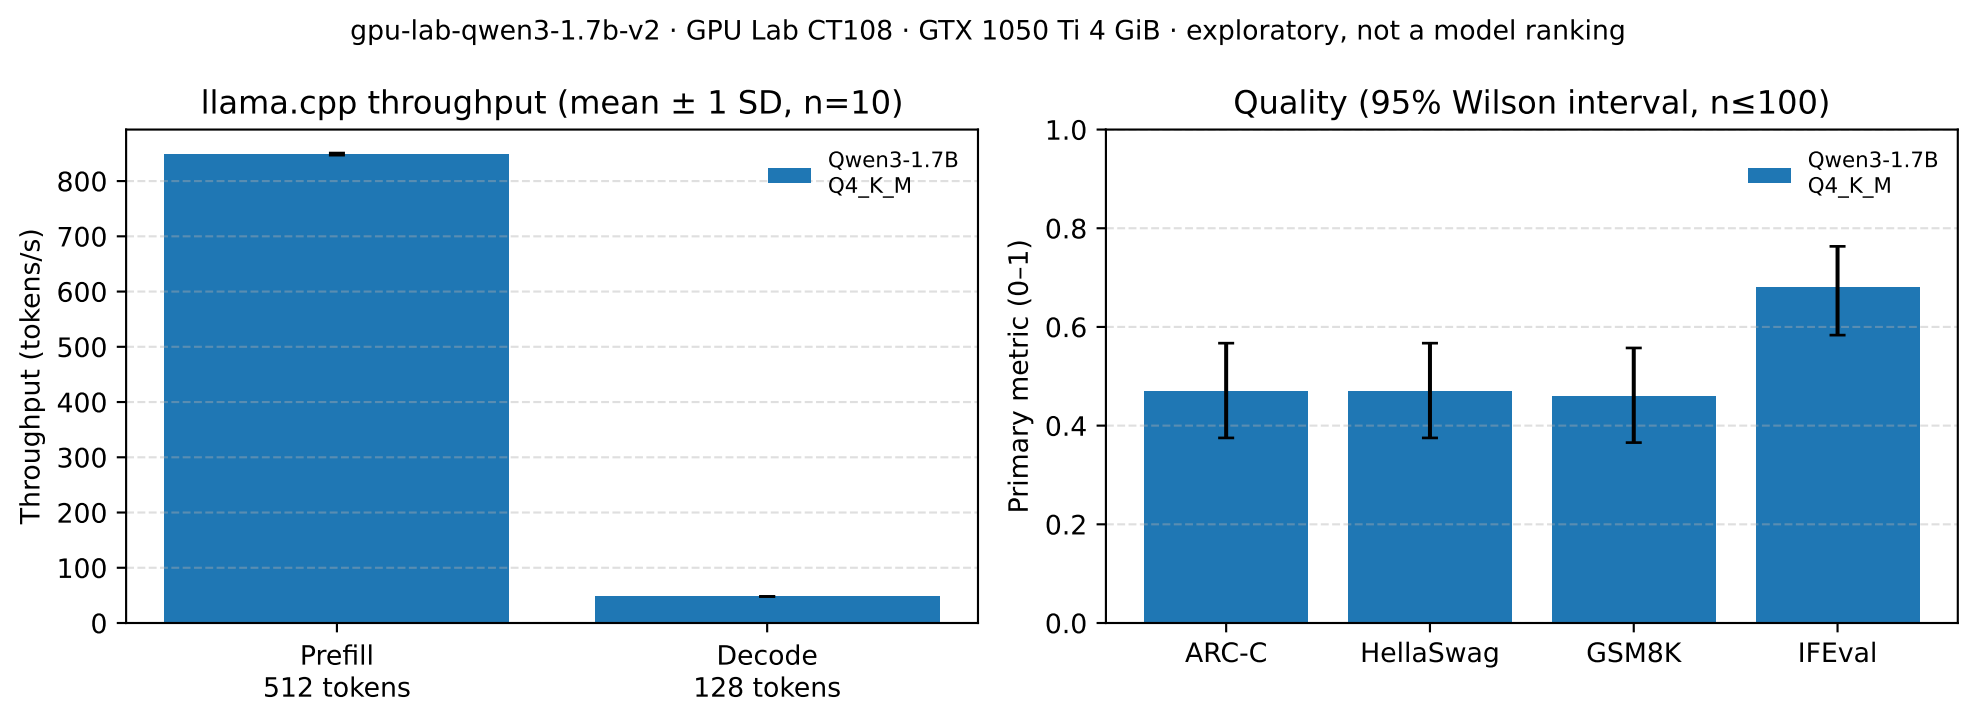
gpu-lab-qwen3-1.7b-v2 · GPU Lab CT108 · GTX 1050 Ti 4 GiB · exploratory, not a model ranking
llama.cpp throughput (mean ± 1 SD, n=10)
Throughput (tokens/s)
0
100
200
300
400
500
600
700
800
Qwen3-1.7B
Q4_K_M
Prefill
512 tokens
Decode
128 tokens
Quality (95% Wilson interval, n≤100)
Primary metric (0–1)
0.0
0.2
0.4
0.6
0.8
1.0
Qwen3-1.7B
Q4_K_M
ARC-C
HellaSwag
GSM8K
IFEval