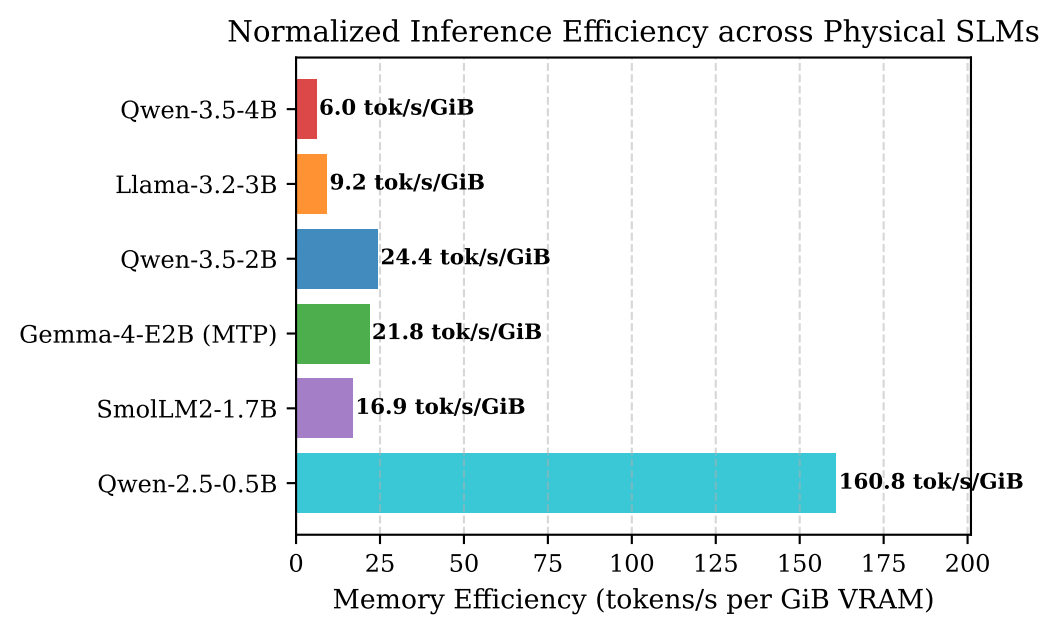
Normalized Inference Efficiency across Physical SLMs
Qwen-3.5-4B
6.0 tok/s/GiB
Llama-3.2-3B
9.2 tok/s/GiB
Qwen-3.5-2B
24.4 tok/s/GiB
Gemma-4-E2B (MTP)
21.8 tok/s/GiB
SmolLM2-1.7B
16.9 tok/s/GiB
Qwen-2.5-0.5B
160.8 tok/s/GiB
0
25
50
75
100
125
150
175
200
Memory Efficiency (tokens/s per GiB VRAM)