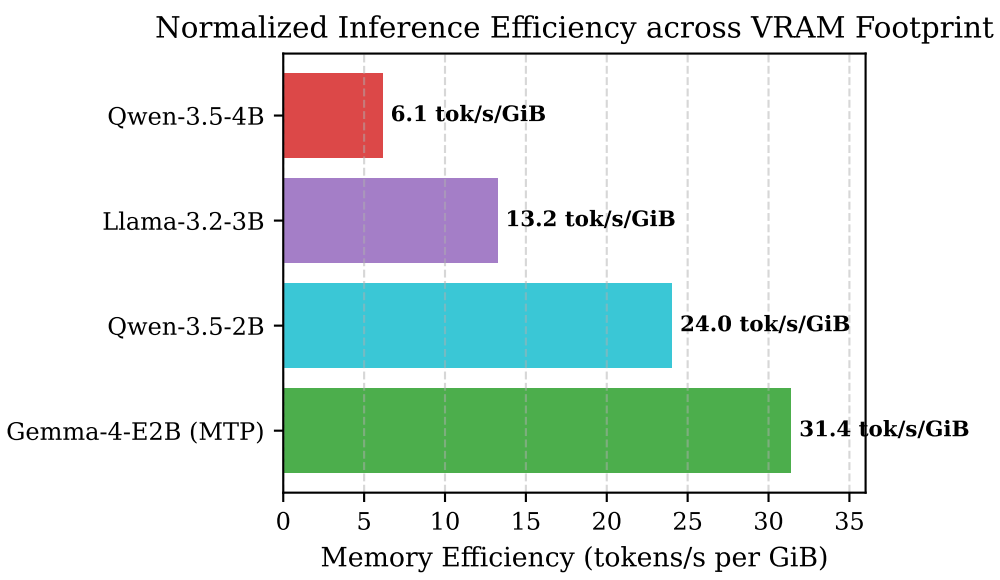
Normalized Inference Efficiency across VRAM Footprint
Qwen-3.5-4B
6.1 tok/s/GiB
Llama-3.2-3B
13.2 tok/s/GiB
Qwen-3.5-2B
24.0 tok/s/GiB
Gemma-4-E2B (MTP)
31.4 tok/s/GiB
0
5
10
15
20
25
30
35
Memory Efficiency (tokens/s per GiB)